M2-BIG DATA
GPGPU - Chapter 8

Exercice 1

Réalisé par:
Gaby Maroun

Encadré par:
Dr. Etancelin JM

February 6, 2021

Objectives

Implement a convolution filter for constant squared boxed filters. The convolution formula is given by :

$$P_{i,j,c} = \sum_{x=-k}^{k} \sum_{y=-k}^{k} I_{i+x,j+y,c} M_{x,y}$$

where c is image channel, M is the filter mask of size $2k+1$ the parameter k is called filter radius.

Instructions

- $k = 1$ and $M = \begin{pmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{pmatrix}$

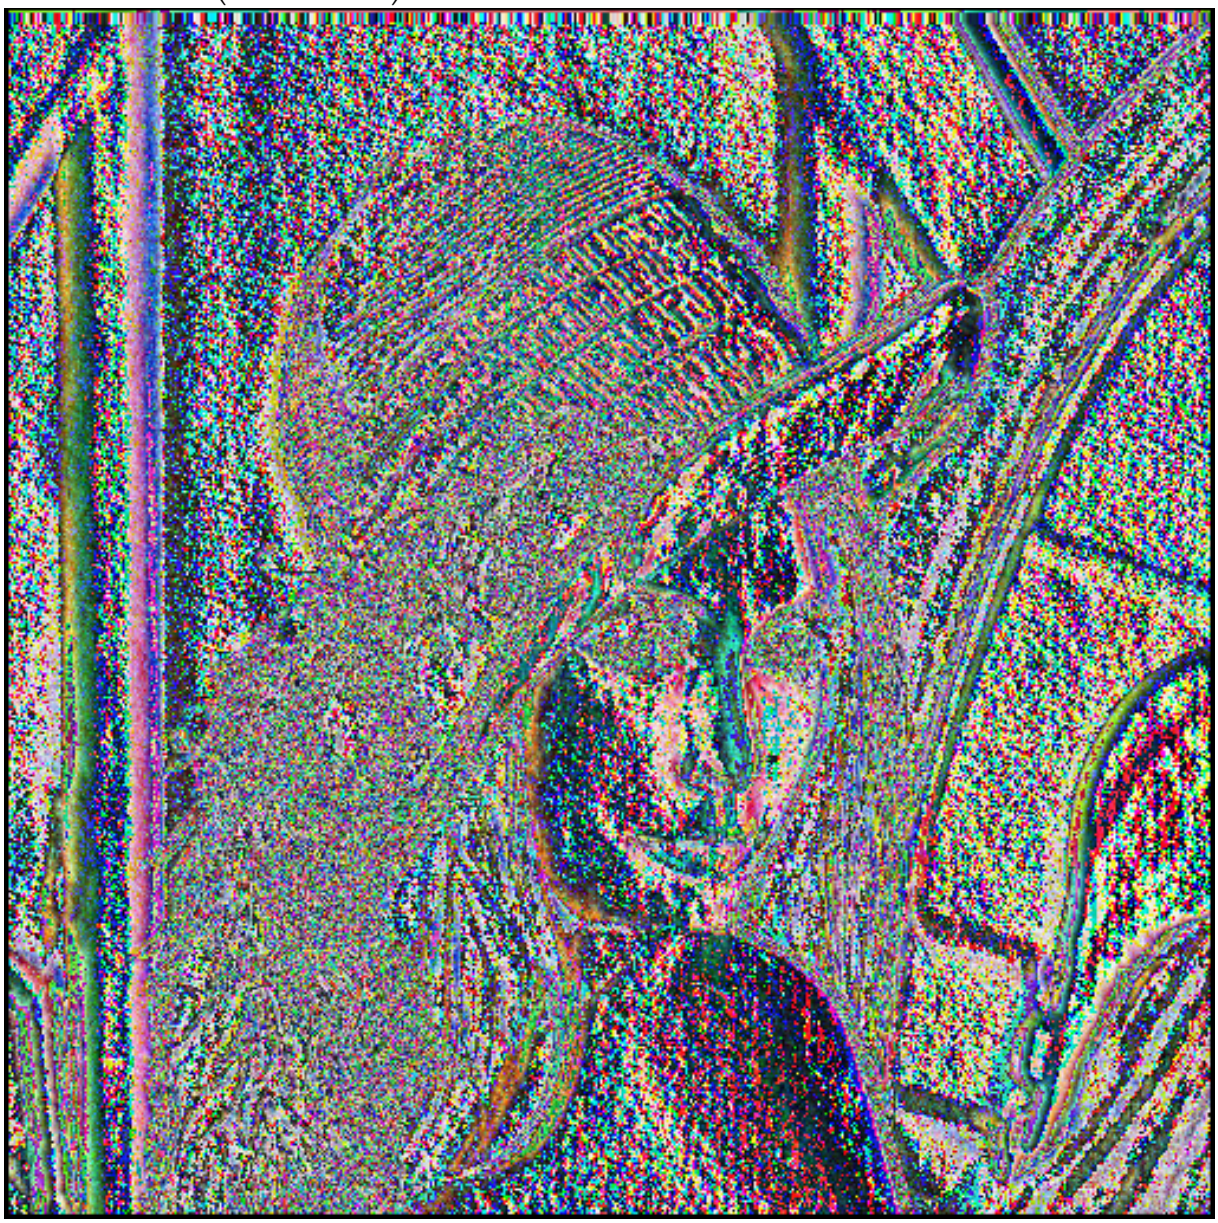

- $k = 1$ and $M = (\frac{1}{25})$

- $k = 1$ and $M = (\frac{-1}{256}) \begin{pmatrix} 1 & 4 & 6 & 4 & 1 \\ 4 & 16 & 24 & 16 & 4 \\ 6 & 24 & -476 & 24 & 6 \\ 4 & 16 & 24 & 16 & 4 \\ 1 & 4 & 6 & 4 & 1 \end{pmatrix}$

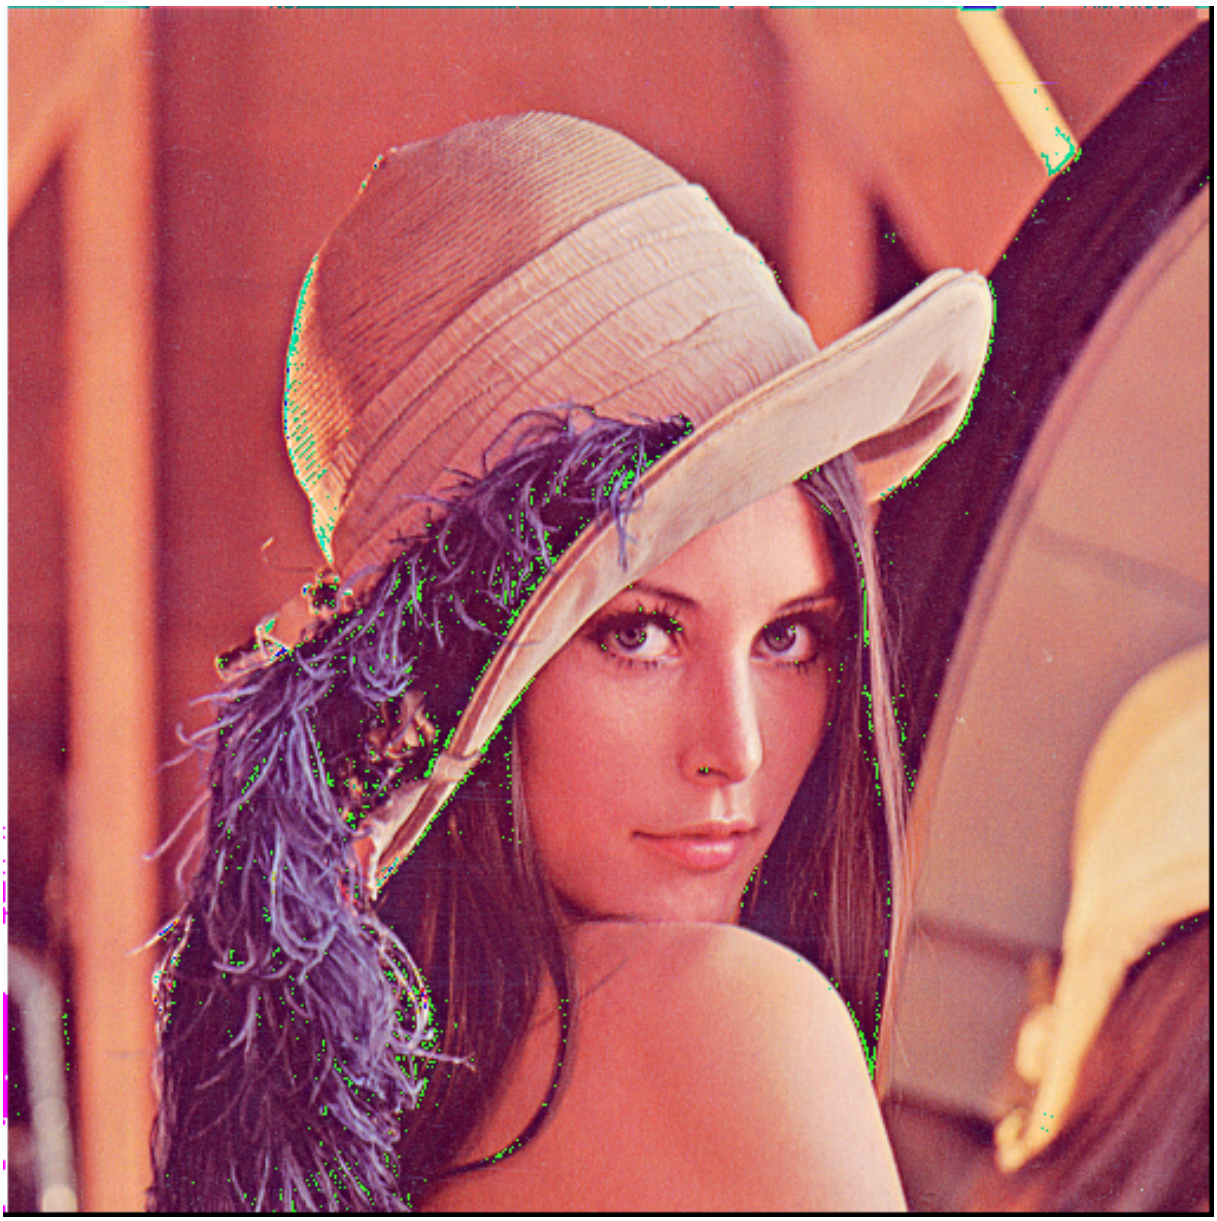

Questions

1. **How many floating operations are being performed in your convolution kernel ? explain.**
   There are

   $$imgCols \times imgRows \times channels(2 \times maskWidth^2)$$

   *floating operations for cycling on all of the image with all the numbers of channels and also cycling through the mask twice, one for + and the other to **

2. **How many global memory reads are being performed by your kernel ? explain.**
   There are

   $$imgCols \times imgRows \times channels(2 \times maskWidth^2)$$

*global memory reads for cycling to read from all the pixels of the image for the different number of channels and also cycling through the mask twice, one for + and the other to **

3. **How many global memory writes are being performed by your kernel ? explain.**
   There are

$$imgCols \times imgRows \times channels$$

   global memory writes for writing the image for different number of channels

4. **Compute the arithmetic intensity of the kernel.**
   The arithmetic intensity is a FLOP/Byte number standing for the number of floating point operations performed per byte of global memory accessed.

$$\frac{imgCols \times imgRows \times channels(2 \times maskWidth^2)}{imgCols \times imgRows \times channels(2 \times maskWidth^2) + imgCols \times imgRows \times channels}$$

$$=> \frac{1}{1 + \frac{imgCols \times imgRows \times channels}{imgCols \times imgRows \times channels(2 \times maskWidth^2)}}(FLOP/Byte)$$

$$=> \frac{1}{1 + \frac{1}{(2 \times maskWidth^2)}}(FLOP/Byte)$$

5. **Measure the kernel computational time of the kernel, using the profiler. Then, compute the computational power of the kernel (in GFLOPS). Compare with the CPU version given.**

   Sequential Version:

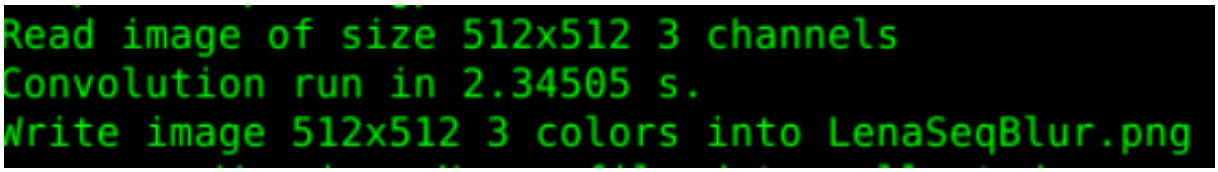

   *As it seems, it takes 2.34505seconds for the kernel to compile with a 25*25 mask on a 512x512 image size.*

   Parallel version:

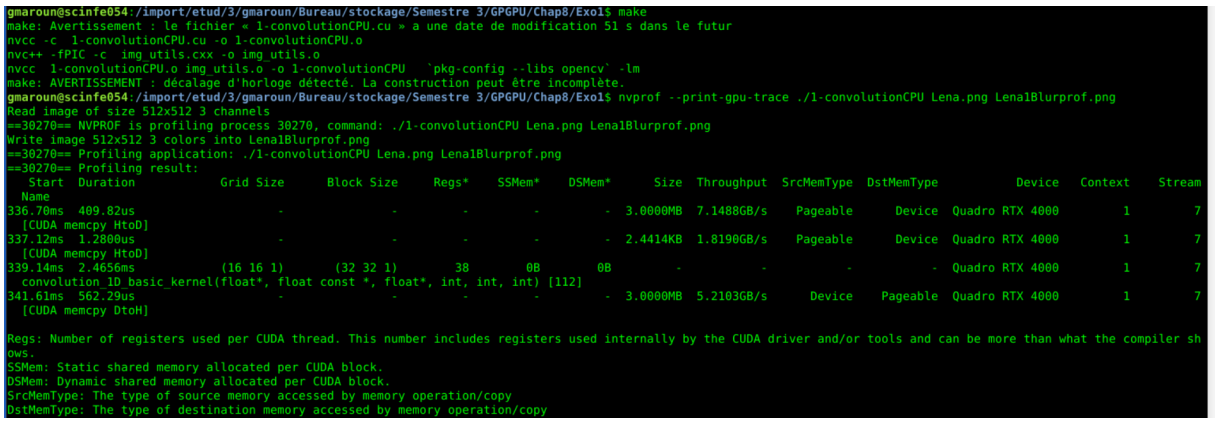

*As it seems, it takes 2.4656milliseconds for the kernel to compile with a 25*25 mask on a 512x512 image size.*

For the computational power of the kernel for the GPU version, it is equal to :

$$\frac{FloatingOperations}{ExecutionTime}$$

$$=> \frac{imgCols \times imgRows \times channels(2 \times maskWidth^2)}{2.4656ms}$$

For the computational power of the kernel for the CPU version and having the same number of floating operations, so :

$$\frac{FloatingOperations}{2.4656 \times 10^3 s} < \frac{FloatingOperations}{2.34505s}$$

That means, this exercise's kernel has 10^3 more computation power than the CPU's

6. **Compare the computational power evolution using different images sizes.**

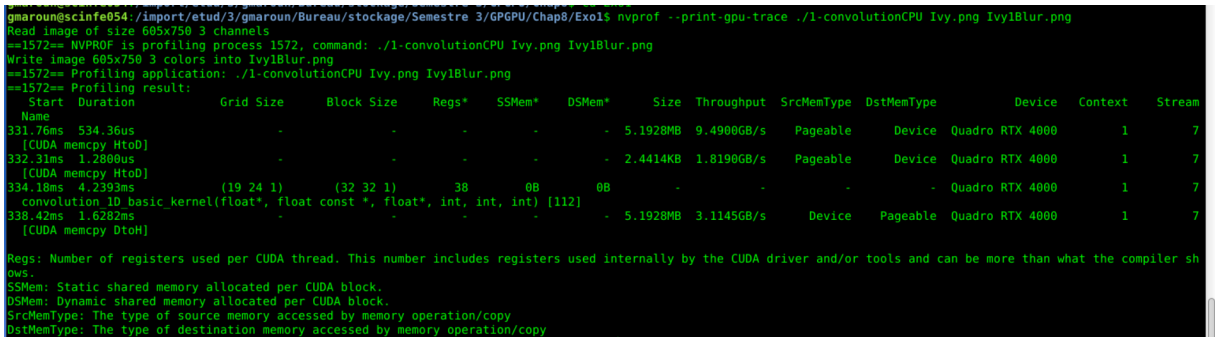

Ivy :

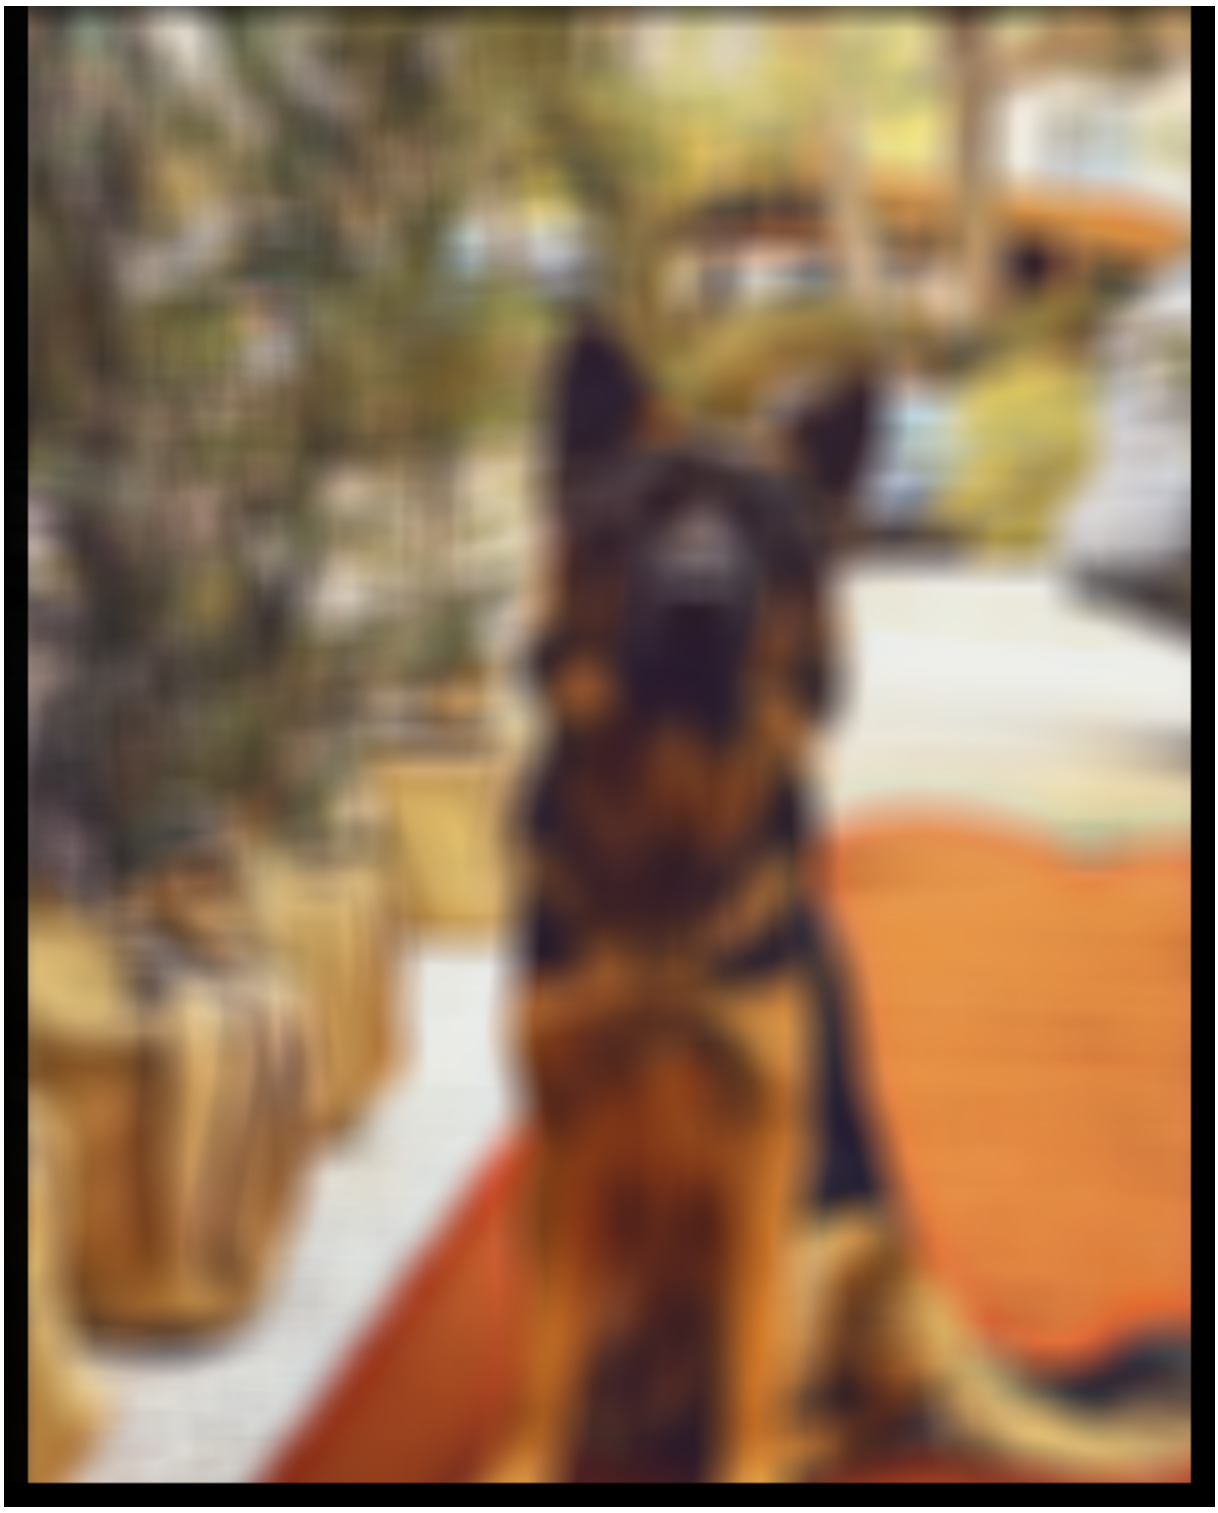

*it takes 4.2393 milliseconds for the kernel to compile with a 25*25 mask on a 605x750 image size*

Tiger4K :

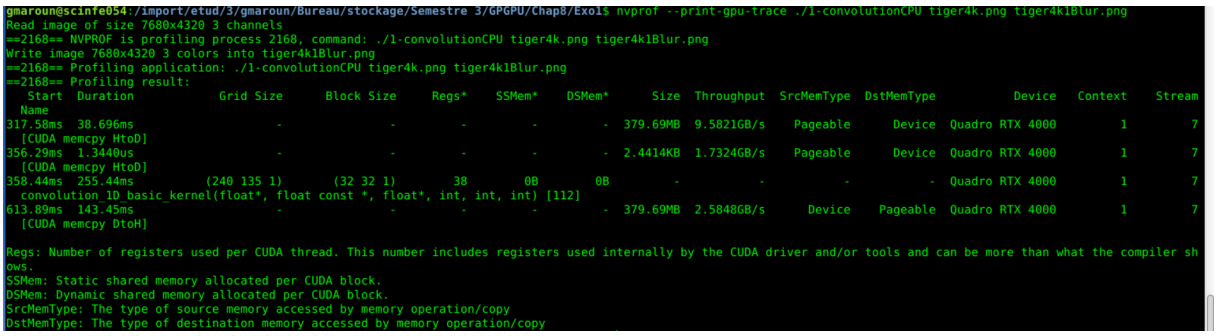

*it takes 255.44milliseconds for the kernel to compile with a 25*25 mask on a 7680x4320 image size*

La fin.