

M2-BIG DATA

GPGPU - Chapter 13

Exercice 1



Réalisé par:
Gaby Maroun

Encadré par:
Dr. Etancelin JM

March 8, 2021

Objectives

Write a first simple OpenACC code.

Instructions

1. Write a CPU version of the vectorAdd code from Chapter 3. *Solution can be found in the file vectorAdd.c*
2. Write a second version using explicit data movement directives *Solution can be found in the file vectorAdd1.c*

Questions

1. Use the profiler to compare the 3 versions of vector add : CUDA, OpenACC with generated data transfers and OpenACC explicit data management. Is the execution time of the whole algorithm (transfers + computations) almost the same ? Explain ?

CPU version of the vectorADD code:

```
gmaroun@scinfo058:/import/etud/3/gmaroun/Bureau/stockage/Semestre 3/GPGPU/Chap13/Exo1$ nvcc -acc-gpu -xopenacc vectorAdd.c
main:
  41, Generating copyin(a[:n],b[:n]) [if not already present]
    Generating copyout(c[:n]) [if not already present]
    Loop is parallelizable
    Generating Tesla code
    41, #pragma acc loop gang, vector(128) /* blockIdx.x threadIdx.x */
gmaroun@scinfo058:/import/etud/3/gmaroun/Bureau/stockage/Semestre 3/GPGPU/Chap13/Exo1$ nvprof --print-gpu-trace ./a.out
Vector size is 1000000
==3142== NVPROF is profiling process 3142, command: ./a.out
Final result:1000000000000.000000
==3142== Profiling application: ./a.out
==3142== Profiling result:
   Start Duration            Grid Size          Block Size      Regs*    SSMem*    DSMem*      Size Throughput  SrcMemType  DstMemType      Device  Context  Stream
   Name
339.22ms  623.00us              -                -              -          -          -  7.6294MB  11.959GB/s   Pinned       Device  Quadro RTX 4000      1       14
[CUDA memcpy HtoD]
340.53ms  621.27us              -                -              -          -          -  7.6294MB  11.993GB/s   Pinned       Device  Quadro RTX 4000      1       14
[CUDA memcpy HtoD]
341.23ms  60.095us             (7813 1 1)        (128 1 1)        16          0B          0B          -          -          -          -          Quadro RTX 4000      1       14
main_41_gpu [34]
341.32ms  629.14us              -                -              -          -          -  7.6294MB  11.842GB/s   Device       Pinned  Quadro RTX 4000      1       14
[CUDA memcpy DtoH]

Regs: Number of registers used per CUDA thread. This number includes registers used internally by the CUDA driver and/or tools and can be more than what the compiler shows.
SSMem: Static shared memory allocated per CUDA block.
DSMem: Dynamic shared memory allocated per CUDA block.
SrcMemType: The type of source memory accessed by memory operation/copy
DstMemType: The type of destination memory accessed by memory operation/copy
```

The second version using explicit data movement directives:

```
gmaroun@scinfo058:/import/etud/3/gmaroun/Bureau/stockage/Semestre 3/GPGPU/Chap13/Exo1$ nvprof --print-gpu-trace ./a.out
Vector size is 1000000
==2866== NVPROF is profiling process 2866, command: ./a.out
Final result:1000000000000.000000
==2866== Profiling application: ./a.out
==2866== Profiling result:
   Start Duration            Grid Size          Block Size      Regs*    SSMem*    DSMem*      Size Throughput  SrcMemType  DstMemType      Device  Context  Stream
   Name
330.45ms  622.26us              -                -              -          -          -  7.6294MB  11.973GB/s   Pinned       Device  Quadro RTX 4000      1       14
[CUDA memcpy HtoD]
331.76ms  621.49us              -                -              -          -          -  7.6294MB  11.988GB/s   Pinned       Device  Quadro RTX 4000      1       14
[CUDA memcpy HtoD]
332.46ms  60.159us             (7813 1 1)        (128 1 1)        16          0B          0B          -          -          -          -          Quadro RTX 4000      1       14
main_44_gpu [34]
332.55ms  629.14us              -                -              -          -          -  7.6294MB  11.842GB/s   Device       Pinned  Quadro RTX 4000      1       14
[CUDA memcpy DtoH]

Regs: Number of registers used per CUDA thread. This number includes registers used internally by the CUDA driver and/or tools and can be more than what the compiler shows.
SSMem: Static shared memory allocated per CUDA block.
DSMem: Dynamic shared memory allocated per CUDA block.
SrcMemType: The type of source memory accessed by memory operation/copy
DstMemType: The type of destination memory accessed by memory operation/copy
```

The CUDA version from chapter 4 exercise 1:

```

gmaroun@scinf058:~/import/etud/3/gmaroun/Bureau/Stockage/Semestre 3/GPGPU/Chap4/Ex1$ nvprof --print-gpu-trace ./1-basicMatMul 1000 1000 1000
Matrix multiplication dimensions: [1000;1000] = [1000;1000] x [1000;1000]
==3561== NVPROF is profiling process 3561, command: ./1-basicMatMul 1000 1000 1000
OK
==3561== Profiling application: ./1-basicMatMul 1000 1000 1000
==3561== Profiling result:
   Start Duration   Grid Size   Block Size   Regs*   SSMem*   DSMem*   Size Throughput SrcMemType DstMemType   Device   Context   Stream
   Name
323.13ms 646.71us      -           -         -         -         - 3.8147MB 5.7604GB/s Pageable Device Quadro RTX 4000 1 7
[CUDA memcpy HtoD]
323.99ms 595.93us      -           -         -         -         - 3.8147MB 6.2513GB/s Pageable Device Quadro RTX 4000 1 7
[CUDA memcpy HtoD]
324.59ms 1.2880us      -           -         -         -         - 3.8147MB 2910.4GB/s Device - Quadro RTX 4000 1 7
[CUDA memset]
326.72ms 4.1438ms (32 32 1) (32 32 1) 48 0B 0B - - - Quadro RTX 4000 1 7
dgemv(float*, float*, float*, int, int, int, int) [113]
330.87ms 1.7872ms      -           -         -         -         - 3.8147MB 2.0844GB/s Device Pageable Quadro RTX 4000 1 7
[CUDA memcpy DtoH]

Regs: Number of registers used per CUDA thread. This number includes registers used internally by the CUDA driver and/or tools and can be more than what the compiler shows.
SSMem: Static shared memory allocated per CUDA block.
DSMem: Dynamic shared memory allocated per CUDA block.
SrcMemType: The type of source memory accessed by memory operation/copy
DstMemType: The type of destination memory accessed by memory operation/copy

```

The execution time of the version is approximately similar by a margin difference of just 20μs. This can be explained by the fact that OpenACC overlaps OpenMP in speed but only work on the accelerators

La fin.